

Data 624_Project 1

Jiaxin Zheng

2025-04-05

Part A

```
library(fpp3)
```

Download the library and load the data

```
## Warning: package 'fpp3' was built under R version 4.4.2

## Registered S3 method overwritten by 'tsibble':
##   method           from
##   as_tibble.grouped_df dplyr

## -- Attaching packages ----- fpp3 1.0.1 --

## v tibble      3.2.1    v tsibble      1.1.6
## v dplyr       1.1.4    v tsibbledata  0.4.1
## v tidyverse    1.3.1    v feasts        0.4.1
## v lubridate    1.9.4    v fable         0.4.1
## v ggplot2     3.5.1

## Warning: package 'dplyr' was built under R version 4.4.3

## Warning: package 'ggplot2' was built under R version 4.4.2

## Warning: package 'tsibbledata' was built under R version 4.4.2

## Warning: package 'feasts' was built under R version 4.4.2

## Warning: package 'fabletools' was built under R version 4.4.2

## Warning: package 'fable' was built under R version 4.4.2

## -- Conflicts ----- fpp3_conflicts --
## x lubridate::date()    masks base::date()
## x dplyr::filter()     masks stats::filter()
## x tsibble::intersect() masks base::intersect()
## x tsibble::interval() masks lubridate::interval()
## x dplyr::lag()        masks stats::lag()
## x tsibble::setdiff()  masks base::setdiff()
## x tsibble::union()    masks base::union()
```

```

library(dplyr)
library(readxl)
library(ggplot2)
library(lubridate)
library(gridExtra)

## 
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
## 
##     combine

library(tsibble)
library(readr)

```

Warning: package 'readr' was built under R version 4.4.3

```
library(openxlsx)
```

Warning: package 'openxlsx' was built under R version 4.4.3

```
atm_data <- read.csv ('https://raw.githubusercontent.com/Jennyjjxxzz/Data-624_Project1/refs/heads/main/atm_data.csv')

head(atm_data)
```

```

##           DATE ATM Cash
## 1 5/1/2009 12:00:00 AM ATM1  96
## 2 5/1/2009 12:00:00 AM ATM2 107
## 3 5/2/2009 12:00:00 AM ATM1  82
## 4 5/2/2009 12:00:00 AM ATM2  89
## 5 5/3/2009 12:00:00 AM ATM1  85
## 6 5/3/2009 12:00:00 AM ATM2  90

```

```

atm_data <- atm_data %>%
  mutate(DATE = mdy_hms(DATE),
        DATE = as.Date(DATE)) %>%
  as_tsibble(index = DATE, key = ATM) %>%
  arrange(DATE)

head(atm_data)
```

Transfer into tsibble dataframe

```

## # A tsibble: 6 x 3 [1D]
## # Key:      ATM [4]
##   DATE      ATM    Cash
```

```

##   <date>      <chr> <int>
## 1 2009-05-01 ATM1     96
## 2 2009-05-01 ATM2    107
## 3 2009-05-01 ATM3     0
## 4 2009-05-01 ATM4    777
## 5 2009-05-02 ATM1     82
## 6 2009-05-02 ATM2    89

```

```

atm_data %>%
  filter(is.na(Cash))

```

Find the missing value from the data.

```

## # A tsibble: 19 x 3 [1D]
## # Key:           ATM [3]
##   DATE      ATM      Cash
##   <date>    <chr> <int>
## 1 2009-06-13 "ATM1"    NA
## 2 2009-06-16 "ATM1"    NA
## 3 2009-06-18 "ATM2"    NA
## 4 2009-06-22 "ATM1"    NA
## 5 2009-06-24 "ATM2"    NA
## 6 2010-05-01 ""       NA
## 7 2010-05-02 ""       NA
## 8 2010-05-03 ""       NA
## 9 2010-05-04 ""       NA
## 10 2010-05-05 ""      NA
## 11 2010-05-06 ""      NA
## 12 2010-05-07 ""      NA
## 13 2010-05-08 ""      NA
## 14 2010-05-09 ""      NA
## 15 2010-05-10 ""      NA
## 16 2010-05-11 ""      NA
## 17 2010-05-12 ""      NA
## 18 2010-05-13 ""      NA
## 19 2010-05-14 ""      NA

```

```

atm_data %>%
  filter_index(~'2010-4-30') %>%
  filter(is.na(Cash))

```

Fillter and find out how many missing data for ATMs

```

## # A tsibble: 5 x 3 [1D]
## # Key:           ATM [2]
##   DATE      ATM      Cash
##   <date>    <chr> <int>
## 1 2009-06-13 ATM1     NA

```

```

## 2 2009-06-16 ATM1      NA
## 3 2009-06-18 ATM2      NA
## 4 2009-06-22 ATM1      NA
## 5 2009-06-24 ATM2      NA

```

```

atm_df <- atm_data %>%
  filter_index(~'2010-04-30') %>%      # Only use data up to April 30, 2010
  fill_gaps() %>%
  filter(!is.na(ATM)) %>%
  group_by(ATM) %>%
  mutate(Cash = if_else(
    is.na(Cash),
    round(mean(Cash, na.rm = TRUE), 0),
    Cash
  )) %>%
  ungroup()

```

```
head(atm_df)
```

Impute the missing values with the means.

```

## # A tsibble: 6 x 3 [1D]
## # Key:      ATM [1]
##   DATE      ATM      Cash
##   <date>    <chr> <dbl>
## 1 2009-05-01 ATM1     96
## 2 2009-05-02 ATM1     82
## 3 2009-05-03 ATM1     85
## 4 2009-05-04 ATM1     90
## 5 2009-05-05 ATM1     99
## 6 2009-05-06 ATM1     88

```

- Double check for missing values

```
sum(is.na(atm_df$Cash))
```

```
## [1] 0
```

Autoplot the atm_data, for better understing for the dataset


- There are outliers for ATM3 and ATM4

```

atm_df %>%
  autoplot(Cash) +
  facet_wrap(~ATM, scale='free_y') +
  ggtitle('ATM Withdrawal for Four ATMs')

```


ATM Withdrawal for Four ATMs



- In the Boxplot, we can see there are some outlier in ATM1 too.

```
plot <- atm_df %>%
  ggplot(aes(x = Cash)) +
  geom_boxplot(aes(fill = ATM)) +
  facet_wrap(~ATM, nrow = 1, scales = "free_x") +
  labs(
    title = "Boxplots of ATM Cash Withdrawals",
    y = "",
    x = "Cash (Hundreds of $)"
  )
plot
```

Boxplots of ATM Cash Withdrawals



Filling the ATMs into the models(ETS, ARIMA, Naive) to see the RMSE

- As a result, we can see that ARIMA is the best model to use for forecasting for all the ATMs because ARIMA has the lowest RMSE.

```
atm_fit_all <- atm_df %>%
  model(
    ETS = ETS(Cash),
    ARIMA = ARIMA(Cash, stepwise = FALSE),
    Naive = NAIVE(Cash),
    snaive = SNAIVE(Cash)
  )

accuracy_results <- atm_fit_all %>%
  accuracy() %>%
  arrange(ATM, RMSE)

accuracy_results

## # A tibble: 16 x 11
##   ATM   .model .type      ME    RMSE     MAE     MPE    MAPE    MASE   RMSSE     ACF1
##   <chr> <chr> <chr>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1 ATM1  ARIMA  Trai~ -7.36e-2  23.4   14.7  -103.   118.  0.825  0.840 -0.00795
## 2 ATM1  ETS    Trai~ -1.80e-1  23.8   15.1  -107.   121.  0.846  0.854  0.125
## 3 ATM1  snaive Trai~ -3.63e-2  27.9   17.8  -96.6  116.  1      1      0.151
```

```

## 4 ATM1 Naive Trai~ -3.02e- 2 50.1 37.7 -132. 167. 2.11 1.80 -0.367
## 5 ATM2 ARIMA Trai~ -8.06e- 1 24.1 17.3 -Inf Inf 0.835 0.803 -0.00893
## 6 ATM2 ETS Trai~ -5.85e- 1 25.0 17.7 -Inf Inf 0.851 0.830 0.0163
## 7 ATM2 snaive Trai~ 2.23e- 2 30.1 20.8 -Inf Inf 1 1 0.0458
## 8 ATM2 Naive Trai~ -4.67e- 2 54.2 42.5 -Inf Inf 2.05 1.80 -0.308
## 9 ATM3 ARIMA Trai~ 2.71e- 1 5.03 0.271 34.6 34.6 0.370 0.625 0.0124
## 10 ATM3 ETS Trai~ 2.70e- 1 5.03 0.273 Inf Inf 0.371 0.625 -0.00736
## 11 ATM3 Naive Trai~ 2.34e- 1 5.09 0.310 28.8 40.2 0.423 0.632 -0.149
## 12 ATM3 snaive Trai~ 7.35e- 1 8.04 0.735 100 100 1 1 0.640
## 13 ATM4 ETS Trai~ -7.89e- 1 647. 310. -556. 588. 0.770 0.722 -0.00587
## 14 ATM4 ARIMA Trai~ -1.21e-10 650. 324. -590. 620. 0.805 0.725 -0.00936
## 15 ATM4 snaive Trai~ -3.70e+ 0 896. 402. -386. 441. 1 1 -0.0150
## 16 ATM4 Naive Trai~ -8.10e- 1 925. 444. -529. 589. 1.10 1.03 -0.488


```

- Plot to see the ACF and PACF plots

```

atm_df %>%
  filter(ATM == "ATM1") %>%
  gg_tsdisplay(Cash, plot_type = 'partial', lag = 30) +
  labs(title = "ATM1 - Time Series, ACF, PACF")

```




```

atm_df %>%
  filter(ATM == "ATM2") %>%
  gg_tsdisplay(Cash, plot_type = 'partial', lag = 30) +
  labs(title = "ATM2 - Time Series, ACF, PACF")


```

ATM2 – Time Series, ACF, PACF




```
atm_df %>%
  filter(ATM == "ATM3") %>%
  gg_tsdisplay(Cash, plot_type = 'partial', lag = 30) +
  labs(title = "ATM3 – Time Series, ACF, PACF")
```

ATM3 – Time Series, ACF, PACF



```
atm_df %>%
  filter(ATM == "ATM4") %>%
  gg_tsdisplay(Cash, plot_type = 'partial', lag = 30) +
  labs(title = "ATM4 - Time Series, ACF, PACF")
```

ATM4 – Time Series, ACF, PACF




Plotting the forecasted data

```
atm_fit <- atm_df %>%
  model(auto_arima = ARIMA(Cash, stepwise = FALSE))

atm_fc <- atm_fit %>%
  forecast(h = "30 days")

atm_fc %>%
  filter(.model == "auto_arima") %>%
  autoplot(atm_df) +
  labs(title = "30-Day Forecast using ARIMA",
       y = "Cash (hundreds)", x = "Date") +
  facet_wrap(~ATM, scales = "free_y")
```

30-Day Forecast using ARIMA



Analysis- Ljung-Box Test:

- The p- value is > 0.05 , that means ARIMA model works well for forecasting.

```
ljung_box_results <- atm_fit %>%
  select(auto_arima) %>%
  augment() %>%
  features(.innov, ljung_box, lag = 10, dof = 3)

ljung_box_results
```

```
## # A tibble: 4 x 4
##   ATM    .model      lb_stat lb_pvalue
##   <chr> <chr>       <dbl>     <dbl>
## 1 ATM1  auto_arima  12.0      0.100
## 2 ATM2  auto_arima   3.78     0.805
## 3 ATM3  auto_arima   0.132     1.00
## 4 ATM4  auto_arima   3.42     0.843
```

```
# Filter forecasts from May 1, 2010
atm_export <- atm_fc %>%
  filter(.model == "auto_arima", DATE >= as.Date("2010-05-01")) %>%
  as_tibble() %>%
  select(ATM, DATE, Forecast = .mean)
```

```

# Create workbook and worksheet
wb <- createWorkbook()
addWorksheet(wb, "ATM Forecasts")
writeData(wb, sheet = "ATM Forecasts", atm_export)

# Save the Excel file
saveWorkbook(wb, "ATM_PartA_30Day_ARIMA_Forecast.xlsx", overwrite = TRUE)

```

Conclusion for part A:

- In part A's analysis, I explored time series forecasting for four different ATMs using cash withdrawal data up to April 30, 2010.
- In the beginning, I prepossessing the data with identifying and imputing five missing values.
- To identify the best forecasting model, I fitted the four models(ETS, Auto ARIME, Naive, and Seasonal Naive) for each ATM. Then, I selected the ARIMA model for all the ATMs for forecast, because ARIMA has lower RMSE than other models.
- Finally, I use residual diagnostics and ljung- box for testing. And the model was satisfied.

Part B

```

power_data <- read.csv ('https://raw.githubusercontent.com/Jennyjjxxzz/Data-624_Project1/refs/heads/main/power_usage.csv')

head(power_data)

```

Load the data

```

##   CaseSequence YYYY.MMM      KWH
## 1           733 1998-Jan 6862583
## 2           734 1998-Feb 5838198
## 3           735 1998-Mar 5420658
## 4           736 1998-Apr 5010364
## 5           737 1998-May 4665377
## 6           738 1998-Jun 6467147

```

- Convert dataframe to tsibble

```

power_data <- power_data %>%
  mutate(Date = yearmonth(YYYY.MMM)) %>%
  as_tsibble(index = Date)

```

```

# There is 1 missing value
sum(is.na(power_data))

```

Check for the missing values.

```

## [1] 1

```

```
power_data %>% filter(if_any(everything(), is.na))
```

```
## # A tsibble: 1 x 4 [1M]
##   CaseSequence YYYY.MMM     KWH      Date
##       <int> <chr>     <int>    <mth>
## 1          861 2008-Sep     NA 2008 Sep
```

Impute the missing value with the mean KWH

```
power_data <- power_data %>%
  mutate(KWH = if_else(
    is.na(KWH),
    round(mean(KWH, na.rm = TRUE)),
    KWH
  ))
```

Check the missing value again

```
sum(is.na(power_data))
```

```
## [1] 0
```


double check the missing value in the row of 861

```
power_data %>%
  filter(CaseSequence==861)
```

```
## # A tsibble: 1 x 4 [1M]
##   CaseSequence YYYY.MMM     KWH      Date
##       <int> <chr>     <dbl>    <mth>
## 1          861 2008-Sep 6502475 2008 Sep
```

```
power_data %>%
  autoplot(KWH) +
  labs(
    title = "Monthly Residential Power Usage (KWH)",
    y = "Kilowatt Hours", x = "Date"
  )
```



Monthly Residential Power Usage (KWH)



Plot the Monthly Residential Power Usage

```
power_data %>%
  gg_tsdisplay(KWH, plot_type='partial') +
  labs(title = "Before Transformation, Residential Power Usage")
```

Before Transformation, Residuals



Plot the Time series, ACF, and PACF(use difference() if need)
Fit and compare each model

- ARIMA model is the best model compare with others, because it has less RMSE.

```
lambda <- power_data %>%
  features(KWH, features = guerrero) %>%
  pull(lambda_guerrero)

power_fit_original <- power_data %>%
  model(
    ARIMA = ARIMA(KWH),
    ETS = ETS(KWH),
    SNaive = SNAIVE(KWH)
  )

accuracy_results_original <- power_fit_original %>%
  accuracy() %>%
  arrange(RMSE)

accuracy_results_original
```

```
## # A tibble: 3 x 10
##   .model .type      ME    RMSE     MAE     MPE    MAPE    MASE   RMSSE   ACF1
##   <chr>  <chr>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1 ARIMA  Training -25069. 827744. 493541. -5.51 11.7 0.708 0.702 0.0128
## 2 ETS    Training  60461. 834738. 505667. -4.37 12.1 0.725 0.708 0.174
```

```

## 3 SNaive Training 94245. 1178792. 697453. -3.94 14.6 1      1    0.231

power_fit_summary <- power_fit_original %>%
  pivot_longer(everything(),
               names_to = "model",
               values_to = "order")

print(power_fit_summary)

## # A mable: 3 x 2
## # Key:   model [3]
##   model                  order
##   <chr>                 <model>
## 1 ARIMA <ARIMA(0,0,1)(1,1,1)[12] w/ drift>
## 2 ETS      <ETS(M,N,M)>
## 3 SNaive    <SNAIVE>

# Using Box-cox transfer the data
lambda <- power_data %>%
  features(KWH, features = guerrero) %>%
  pull(lambda_guerrero)

power_data <- power_data %>%
  mutate(KWH_trans_diff = difference(box_cox(KWH, lambda)))


power_data %>%
  gg_tsdisplay(KWH_trans_diff, plot_type = "partial", lag = 12) +
  ggtitle("Box-Cox Transformed and Differenced KWH")

## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_line()').

## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_point()').

```

Box–Cox Transformed and Differenced KWH




```
power_arima_fit <- power_data %>%
  model(ARIMA = ARIMA(KWH))

power_arima_fc <- power_arima_fit %>%
  forecast(h = "12 months")

power_arima_fc %>%
  autoplot(power_data) +
  labs(
    title = "Forecast of Residential Power Usage",
    y = "KWH",
    x = "Date"
  )
```

Forecast of Residential Power



Forecast of power usage by using ARIMA model for 12 months

```
report(power_arima_fit)
```

```
## Series: KWH
## Model: ARIMA(0,0,1)(1,1,1)[12] w/ drift
##
## Coefficients:
##          ma1      sar1      sma1    constant
##          0.2443   -0.1520   -0.7309  105249.87
## s.e.  0.0723    0.0963    0.0821   27038.86
##
## sigma^2 estimated as 7.474e+11: log likelihood=-2720
## AIC=5450    AICc=5450.34    BIC=5465.96
```

- The model works well, the P-values are > 0.05

```
power_arima_fit %>%
  augment() %>%
  features(.innov, ljung_box, lag = 10, dof = 2)
```

```
## # A tibble: 1 x 3
##   .model lb_stat lb_pvalue
##   <chr>     <dbl>     <dbl>
## 1 ARIMA     7.28     0.507
```

```

# convert to data frame
power_export <- power_arima_fc %>%
  filter(year(Date) == 2014) %>%
  as_tibble() %>%
  select(Date, Forecast = .mean)

# Create a new Excel workbook and add sheet
wb <- createWorkbook()
addWorksheet(wb, "Power Forecast 2014")
writeData(wb, sheet = "Power Forecast 2014", power_export)

# Save Excel file
saveWorkbook(wb, "PartB_Power_Forecast_2014.xlsx", overwrite = TRUE)

```

Conclusion for part B:

- In part B's analysis, I analyzed the monthly residential power data from January 1988 to December 2013, and forecasted the usage for the year of 2014. The original data show the data is non-stationary. To stabilize, I applied the Box-cox transformation. I evaluated the models, and ARIMA is the best model for this data. At the end, I used ljung- box for the test, and the model was satisfied.

Part C – BONUS

```

pipe1 <- read.csv ('https://raw.githubusercontent.com/Jennyjjxxzz/Data-624_Project1/refs/heads/main/Wat
pipe2 <- read.csv ('https://raw.githubusercontent.com/Jennyjjxxzz/Data-624_Project1/refs/heads/main/Wat

str(pipe1)

```

Load the datas

```

## 'data.frame':    1000 obs. of  2 variables:
##   $ Date.Time: chr  "10/23/15 12:24 AM" "10/23/15 12:40 AM" "10/23/15 12:53 AM" "10/23/15 12:55 AM" ...
##   $ WaterFlow: num  23.4 28 23.1 30 6 ...

str(pipe2)

```

```

## 'data.frame':    1000 obs. of  2 variables:
##   $ Date.Time: chr  "10/23/15 1:00 AM" "10/23/15 2:00 AM" "10/23/15 3:00 AM" "10/23/15 4:00 AM" ...
##   $ WaterFlow: num  18.8 43.1 38 36.1 31.9 ...

```

Aggregate Data

- Transform the Date.Time columns to actual date types.

```

pipe1 <- pipe1 %>%
  mutate(Date.Time = as.POSIXct(Date.Time, format = "%m/%d/%y %I:%M %p"))

pipe2 <- pipe2 %>%
  mutate(Date.Time = as.POSIXct(Date.Time, format = "%m/%d/%y %I:%M %p"))

```

```

pipe_df <- rbind(pipe1, pipe2)

head(pipe_df)

```

To combine the two data into one:

```

##           Date.Time WaterFlow
## 1 2015-10-23 00:24:00 23.369599
## 2 2015-10-23 00:40:00 28.002881
## 3 2015-10-23 00:53:00 23.065895
## 4 2015-10-23 00:55:00 29.972809
## 5 2015-10-23 01:19:00  5.997953
## 6 2015-10-23 01:23:00 15.935223

```

```

pipe_df %>% filter(if_any(everything(), is.na))

```

Check for missing values in data

```

## [1] Date.Time WaterFlow
## <0 rows> (or 0-length row.names)

```

- Now take an average for each hour.

```

pipe_df <- pipe_df %>%
  mutate(Date.Time = floor_date(Date.Time, "hour")) %>%
  group_by(Date.Time) %>%
  summarize(avg_flow = mean(WaterFlow, na.rm=T)) %>%
  as_tsibble(index=Date.Time)

head(pipe_df)

```


```

## # A tsibble: 6 x 2 [1h] <?>
##   Date.Time           avg_flow
##   <dttm>              <dbl>
## 1 2015-10-23 00:00:00    26.1
## 2 2015-10-23 01:00:00    18.8
## 3 2015-10-23 02:00:00    24.5
## 4 2015-10-23 03:00:00    25.6
## 5 2015-10-23 04:00:00    22.4
## 6 2015-10-23 05:00:00    23.6

```


- Plot the data

```
pipe_df %>%
  autoplot(avg_flow) +
  ggtitle('Average waterflow vs Date.hour')
```



- We can see in ACF plots, the data is non-stationary, so we have to transfer

```
pipe_df %>%
  gg_tsdisplay(avg_flow, plot_type = 'partial', lag = 24) + ggtitle('Before Transfer')
```



```
#find lambad
lambad_pipe <- pipe_df %>%
  features(avg_flow, features = guerrero) %>%
  pull (lambda_guerrero)

#find ndiffs
pipe_df %>%
  mutate(avg_flow = box_cox(avg_flow, lambad_pipe)) %>%
  features(avg_flow, unitroot_ndiffs)
```


Transfer the data

```
## # A tibble: 1 x 1
##   ndiffs
##   <int>
## 1     1

pipe_df %>%
  gg_tsdisplay(difference(box_cox(avg_flow, lambda = lambad_pipe)), plot_type='partial', lag = 24) +
  labs(title = paste("Differenced average waterflow"))

## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_line()').
```

```
## Warning: Removed 1 row containing missing values or values outside the scale range
##   ('geom_point()').
```




```
pipe_fit_arima <- pipe_df %>%
  model(ARIMA = ARIMA(box_cox(avg_flow, lambda = lambad_pipe)))
```

```
pipe_forecast_arima <- pipe_fit_arima %>%
  forecast(h = "168 hours") # a week

pipe_forecast_arima %>%
  autoplot(pipe_df) +
  labs(
    title = "1-Week Forecast of Average Water Flow (ARIMA with Box-Cox)",
    x = "Time (Hourly)",
    y = "Average Flow"
  )
```

1-Week Forecast of Average Water Flow (ARIMA with Box-Cox)



Forecast the data

- The p-value > 0.05

```
ljung_box_pipe %>%  
  augment() %>%  
  features(.innov, ljung_box, lag = 24, dof = 2)  
  
ljung_box_pipe  
  
## # A tibble: 1 x 3  
##   .model lb_stat lb_pvalue  
##   <chr>    <dbl>     <dbl>  
## 1 ARIMA     29.0      0.146  
  
pipe_export %>% pipe_forecast_arima %>%  
  as_tibble() %>%  
  select(DateTime = Date.Time, Forecast = .mean)  
  
# Create a new workbook and write to Excel  
wb <- createWorkbook()  
addWorksheet(wb, "Pipe Forecast")  
writeData(wb, sheet = "Pipe Forecast", pipe_export)  
  
# Save the Excel file for Part C  
saveWorkbook(wb, "PartC_Pipe1_1Week_Forecast.xlsx", overwrite = TRUE)
```

Conclusion for part C:

- In the part C, I worked with two data sets, and combined as one. They representing water flow rates at different timestamps for two different pipelines. The goal is to forecast waterflow one week ahead. In this data, the data is non-stationary, and I use box-cox to transfer the data. I selected the ARIMA model for final forecasting. Finally, I use Ljung- Box for test, and the result show the model ARIMA works well, the residual resembled white noise.